



Data File Partitioning

Prof. Hyuk-Yoon Kwon

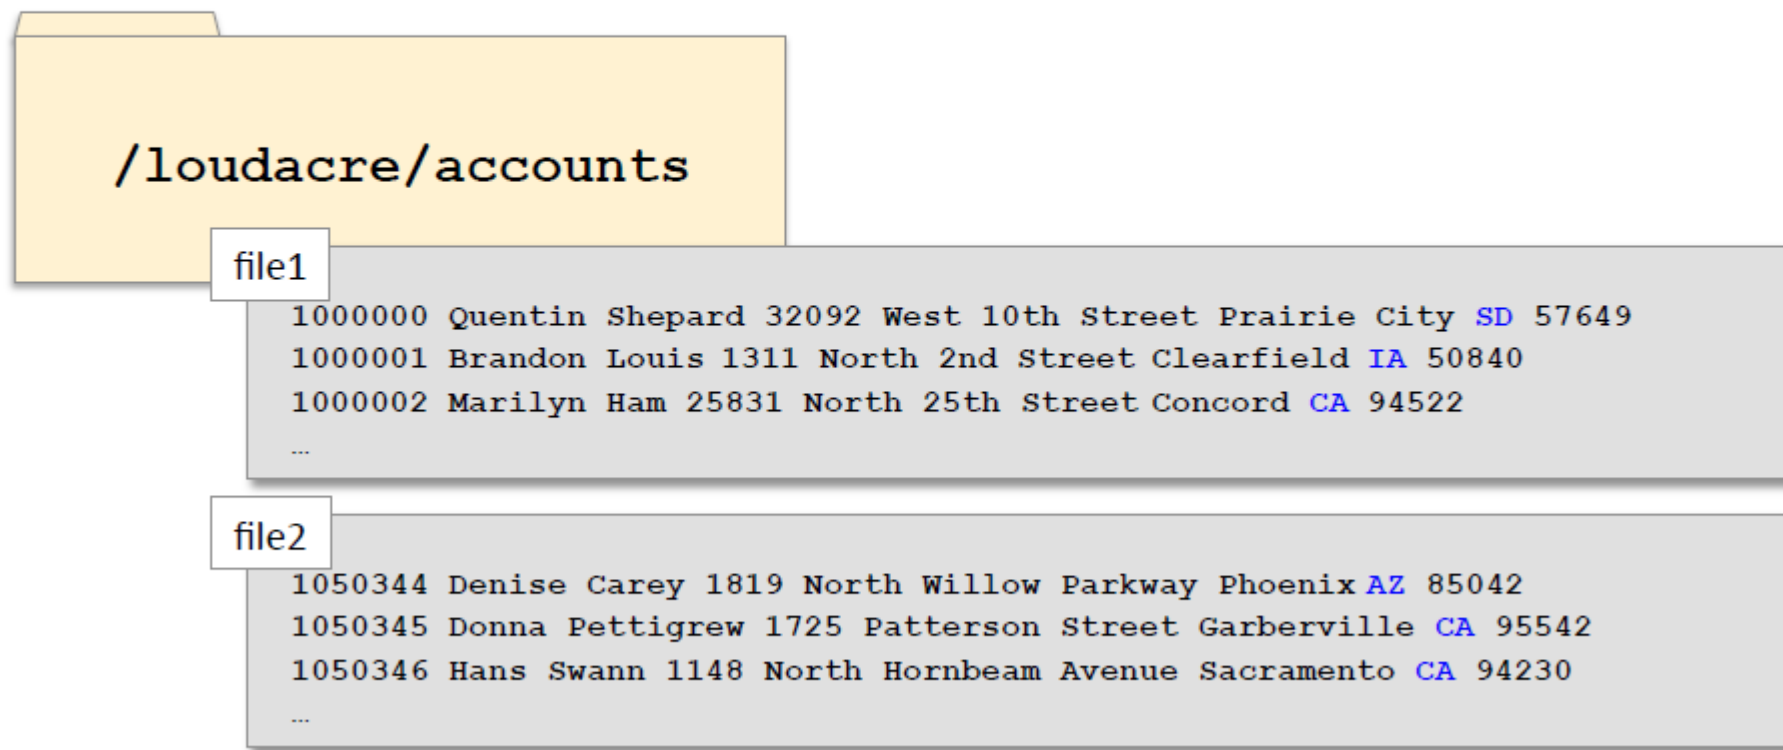
Data File Partitioning

In this chapter you will learn

- How to improve query performance with data file partitioning
- How to create and populate partitioned tables in Impala and Hive

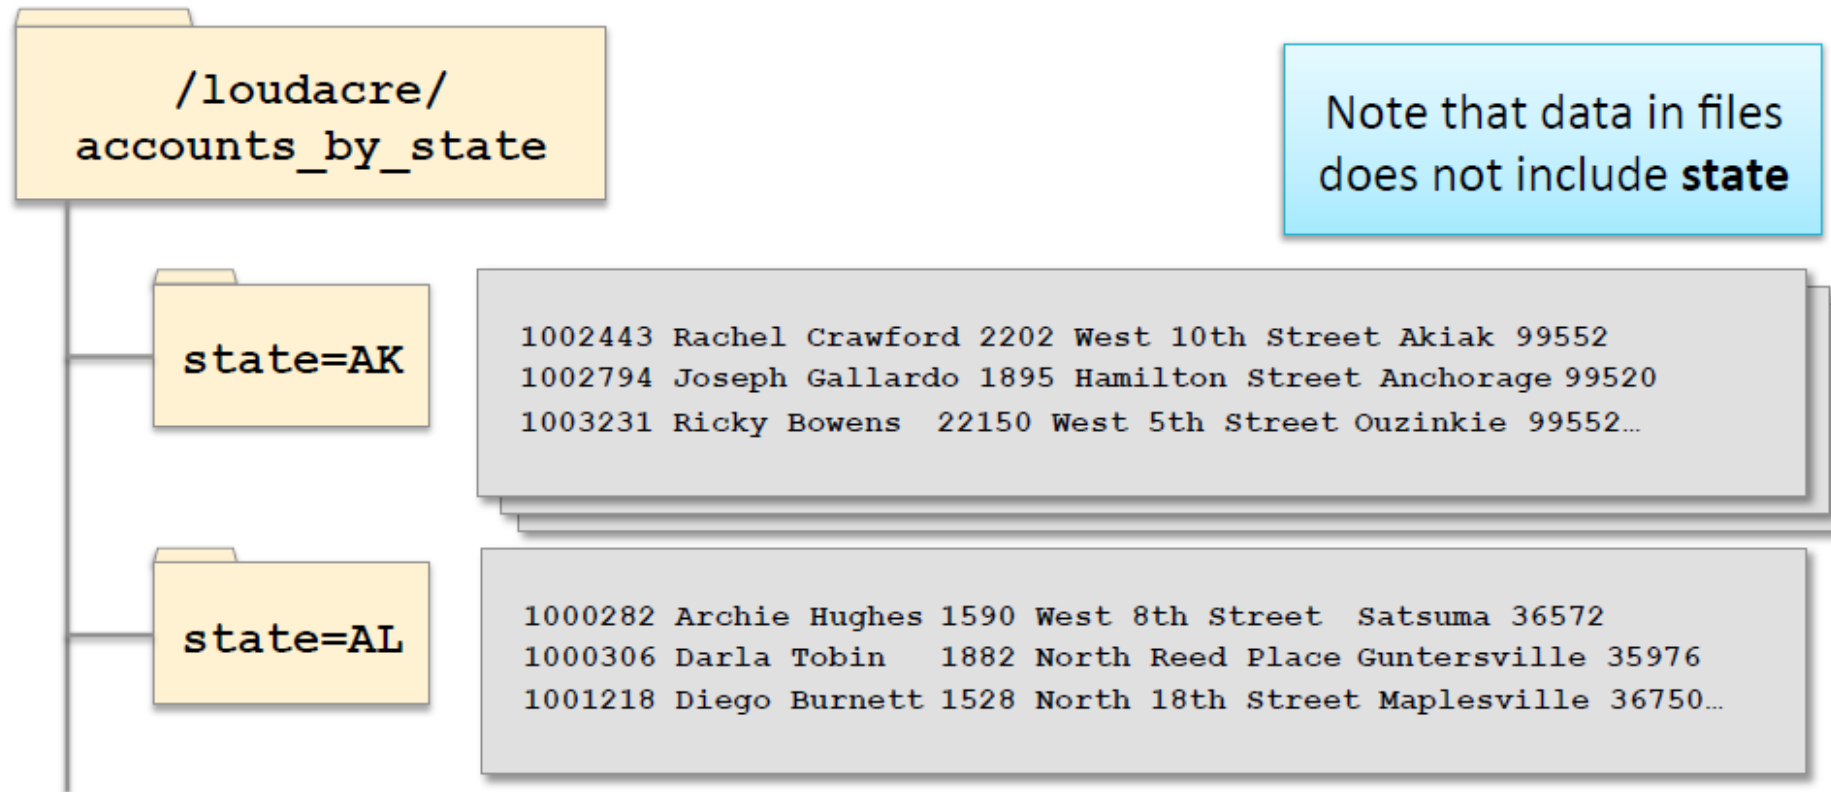
Data Storage Partitioning (1)

- By default, all files in a data set are stored in a single HDFS directory
 - All files in the directory are read during analysis or processing
 - “Full table scan”



Data Storage Partitioning (2)

- *Partitioning* subdivides the data
 - Analysis can be done on only the relevant subset of data
 - Potentially much faster!
- Hadoop partitions using subdirectories



Hadoop Partitioning

- **Partitioning is involved at two phases**
 - Storage – putting the data into correct partition (subdirectory)
 - Retrieval – getting the data out of the correct partition based on the query or analysis being done
- **Hadoop with built-in support for partitioning**
 - Hive and Impala (covered in next section)
 - Sqoop – When using the `--hive-import` option you can specify flags `--hive-partition-key` and `--hive-partition-value`
- **Other tools can be used to store partitioned data**
 - Spark and MapReduce
 - Flume (at ingestion)

Example: Impala/Hive Partitioning Accounts By State (1)

- Example: `accounts` is a non-partitioned table

```
CREATE EXTERNAL TABLE accounts(  
    cust_id INT,  
    fname STRING,  
    lname STRING,  
    address STRING,  
    city STRING,  
    state STRING,  
    zipcode STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
LOCATION '/loudacre/accounts';
```

Example: Impala/Hive Partitioning Accounts By State (2)

- What if most of Loudacre's analysis on the customer table was done by state? For example:

```
SELECT fname, lname  
FROM accounts  
WHERE state='NY';
```

- By default, all queries have to scan *all* files in the directory
- Use partitioning to store data in separate files by state
 - State-based queries scan only the relevant files

Example: Impala/Hive Partitioning Accounts By State (3)

- Create a partitioned table using PARTITIONED BY

```
CREATE EXTERNAL TABLE accounts_by_state(  
    cust_id INT,  
    fname STRING,  
    lname STRING,  
    address STRING,  
    city STRING,  
    state STRING,  
    zipcode STRING)  
PARTITIONED BY (state STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
LOCATION '/loudacre/accounts_by_state';
```


Partition Columns

- The partition column is displayed if you DESCRIBE the table

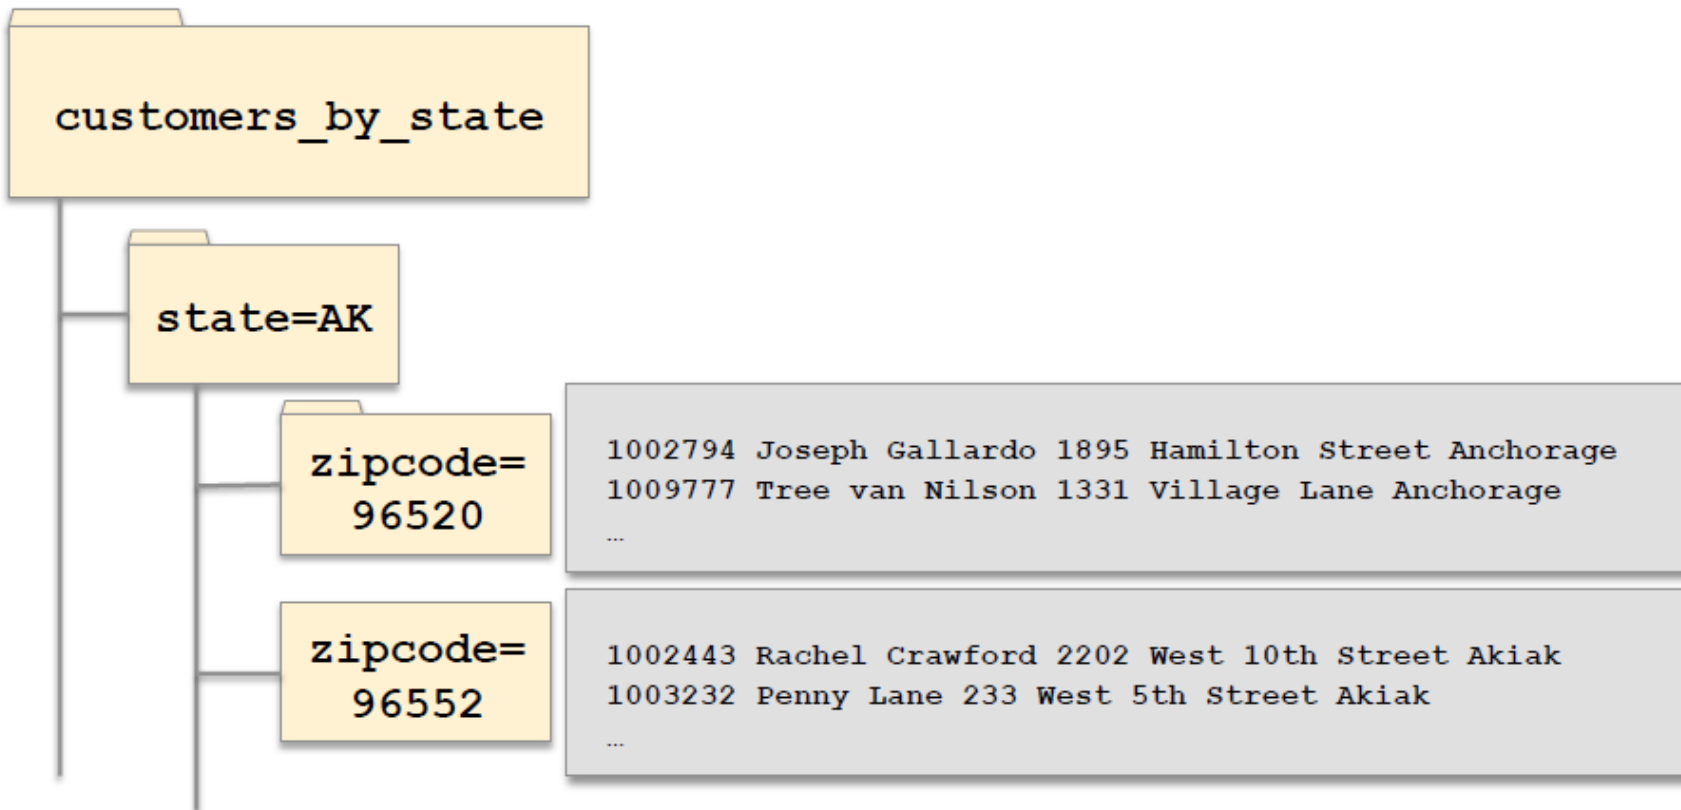
```
DESCRIBE accounts_by_state;  
+-----+-----+-----+  
| name   | type   | comment |  
+-----+-----+-----+  
| cust_id | int    |         |  
| fname   | string |         |  
| lname   | string |         |  
| address | string |         |  
| city    | string |         |  
| zipcode | string |         |  
| state   | string |         |  
+-----+-----+-----+
```

A partition column is a “virtual column”; data is not stored in the file

Nested Partitions

- You can also create nested partitions

```
... PARTITIONED BY (state STRING, zipcode STRING)
```



Loading Data Into a Partitioned Table

- **Dynamic partitioning**

- Impala/Hive add new partitions automatically as needed at load time
- Data is stored into the correct partition (subdirectory) based on column value

- **Static partitioning**

- You define new partitions using `ADD PARTITION`
- When loading data, you specify which partition to store data in

Dynamic Partitioning

- We can create new partitions dynamically from existing data

```
INSERT OVERWRITE TABLE accounts_by_state
  PARTITION(state)
  SELECT cust_id, fname, lname, address,
         city, zipcode, state FROM accounts;
```

- Partitions are automatically created based on the value of the *last* column
 - If the partition does not already exist, it will be created
 - If the partition does exist, it will be overwritten

Practice – Dynamic Partitioning

1. Import the **accounts** table from MySQL directly into the Hive
 - Hint1: use Sqoop command
 - Hint2: use the option “--hive-import”
2. Create a partitioned table by the **city** from **accounts** table
3. Insert data into the partitioned table from **accounts**
4. Create a partitioned table **accounts_areacode** by the **areacode** from **accounts** table
 - areacode could be defined by the first three digits of the phone number
5. Insert data into the partitioned table **accounts_areacode** from **accounts**
 - areacode could be obtained by substr(phone_number, 1, 3)
6. Create a nested partitioned table **accounts_nested** by (state, areacode) from accounts table
7. Insert data into the nested partitioned table **accounts_nested** from **accounts**

Static Partitioning Example: Partition Calls by Day (1)

- Loudacre's customer service phone system generates logs detailing calls received
 - Analysts use this data to summarize previous days' calls
 - For example:

```
SELECT event_type, COUNT(event_type)
FROM call_log
WHERE call_date = '2014-10-01'
GROUP BY event_type;
```

Static Partitioning Example: Partition Calls by Day (2)

- Logs are generated daily, e.g.

call-20141001.log

```
19:45:19,312-555-7834,CALL_RECEIVED
19:45:23,312-555-7834,OPTION_SELECTED,Shipping
19:46:23,312-555-7834,ON_HOLD
19:47:51,312-555-7834,AGENT_ANSWER,Agent ID N7501
19:48:37,312-555-7834,COMPLAINT,Item not received
19:48:41,312-555-7834,CALL_END,Duration: 3:22
...
```

call-20141002.log

```
03:45:01,505-555-2345,CALL_RECEIVED
03:45:09,505-555-2345,OPTION_SELECTED,Billing
03:56:21,505-555-2345,AGENT_ANSWER,Agent ID A1503
03:57:01,505-555-2345,QUESTION
...
```

Static Partitioning Example: Partition Calls by Day (3)

- In the previous example, existing data was partitioned dynamically based on a column value
- This time we use static partitioning
 - Because the data files do not include the partitioning data

Static Partitioning Example: Partition Calls by Day (4)

- The partitioned table is defined the same way

```
CREATE TABLE call_logs (  
    call_time STRING,  
    phone STRING,  
    event_type STRING,  
    details STRING)  
PARTITIONED BY (call_date STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ',';
```

Loading Data Into Static Partitions (1)

- With static partitioning, you create new partitions as needed
- e.g. For each new day of call log data, add a partition:

```
ALTER TABLE call_logs  
  ADD PARTITION (call_date='2014-10-02');
```

- **This command**
 1. Adds the partition to the table's metadata
 2. Creates subdirectory
/user/hive/warehouse/call_logs/
call_date=2014-10-02

Loading Data Into Static Partitions (2)

- Then load the day's data into the correct partition

```
LOAD DATA INPATH '/mystaging/call-20141002.log'  
  INTO TABLE call_logs  
  PARTITION(call_date='2014-10-02');
```

- This command moves the HDFS file `call-20141002.log` to the partition subdirectory

- To overwrite all data in a partition

```
LOAD DATA INPATH '/mystaging/call-20141002.log'  
  INTO TABLE call_logs OVERWRITE  
  PARTITION(call_date='2014-10-02');
```

Practice – Static Partitioning

1. Create the following partitioned table

```
create external table temp(id int, name string, sal int)
partitioned by(city string)
location '/loudacre/test';
```

2. Make static partitions for city='hyd' and city='sec' using alter table

3. Make two sample text files that contains the following contents in HDFS

[test1.txt]

1 ravi 100 hyd

2 krishna 200 hyd

[test2.txt]

3 fff 300 sec

4. Load two sample files to the table for the created static partitions